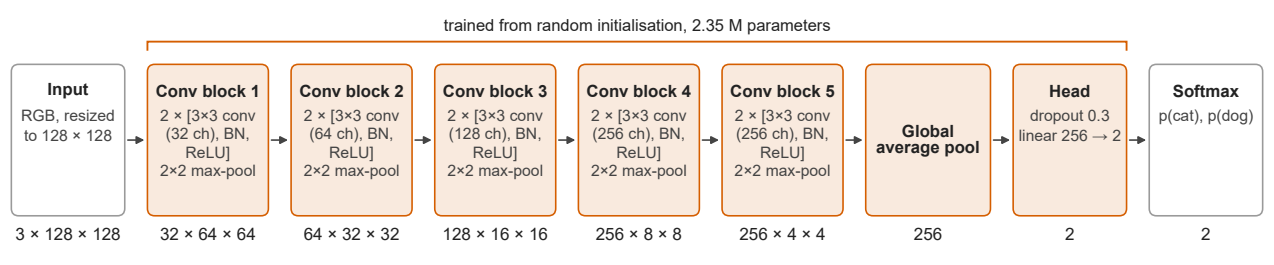
trained from random initialisation, 2.35 M parameters
Input
RGB, resized to 128 × 128
Conv block 1
2 × [3×3 conv (32 ch), BN, ReLU]
2×2 max-pool
Conv block 2
2 × [3×3 conv (64 ch), BN, ReLU]
2×2 max-pool
Conv block 3
2 × [3×3 conv (128 ch), BN, ReLU]
2×2 max-pool
Conv block 4
2 × [3×3 conv (256 ch), BN, ReLU]
2×2 max-pool
Conv block 5
2 × [3×3 conv (256 ch), BN, ReLU]
2×2 max-pool
Global average pool
Head
dropout 0.3
linear 256 → 2
Softmax
p(cat), p(dog)
3 × 128 × 128
32 × 64 × 64
64 × 32 × 32
128 × 16 × 16
256 × 8 × 8
256 × 4 × 4
256
2
2